



PW SKILLS | DevOps and Cloud Computing

Terraform for Cost Management, Security, and AWS Infrastructure Automation



Objective

- Implement cost-saving strategies in cloud infrastructure using Terraform.
- Secure Terraform configurations, credentials, and state files using encryption and IAM policies.
- Automate AWS infrastructure including EC2, S3, VPCs, and RDS using Terraform.
- Configure auto-scaling, monitoring, and backup solutions in AWS with Terraform.





Explaining key cost optimization strategies in cloud infrastructure:

Let's see

1. Reserved & Spot Instances:

- Reserved Instances offer significant discounts for long-term commitments.
- Spot Instances provide low-cost, short-term compute capacity for flexible or non-critical workloads.

2. Auto-Scaling & Right-Sizing:

- Auto-scaling adjusts resources automatically based on demand, preventing overuse.
- Right-sizing matches instance types/sizes to workload needs, avoiding over-provisioning.

3. Data Optimization:

- Use compression to reduce storage size.
- Move infrequently accessed data to archival storage like AWS Glacier or Azure Archive for lower costs.

Pop Quiz

Q. Spot Instances are best suited for which of the following scenarios?

A

Workloads that can tolerate
interruptions

B

Predictable and stable workloads

Pop Quiz

Q. Spot Instances are best suited for which of the following scenarios?

A

Workloads that can tolerate
interruptions

B

Predictable and stable workloads



**Discussing how
Terraform helps enforce
cost controls through
policies and automated
scaling.**

Let's discuss

Terraform helps enforce cost controls by:

- **Policies with Sentinel:** Enforces rules (e.g., restrict instance types, limit regions) to prevent costly configurations.
- **Automated Scaling:** Manages infrastructure as code, enabling auto-scaling groups and resource adjustments based on demand.
- **Consistent Provisioning:** Avoids manual errors and ensures cost-efficient setups across environments.





**Explaining best practices
for version control of
Terraform configurations
using Git.**

Let's see

Best practices for version control of Terraform configs using Git:

- Use Git Repos
- Branching Strategy
- Use .gitignore
- Code Reviews
- Tag Releases
- Commit Often





**Discussing branching
strategies for Terraform
workflows (feature
branches, main/stable
environments).**

Let's discuss

Branching strategies for Terraform workflows:

- **Feature Branches:** Used for developing and testing changes safely without affecting main infrastructure.
- **Main Branch:** Holds stable, reviewed, and approved code—typically reflects production-ready infrastructure.
- **Environment Branches:** Separate branches (e.g., dev, staging, prod) represent different deployment environments for isolated changes and testing.
- **Pull Requests:** Used to merge feature branches into main/stable branches after review and validation.





**Demonstrating how to
structure Terraform code
in a Git repository for
modular reuse.**

Let's do it

To structure Terraform code in a Git repo for modular reuse:

- **modules/ Folder:** Store reusable modules (e.g., modules/vpc, modules/ec2).
- **environments/ Folder:** Separate configs for each environment (e.g., environments/dev, environments/prod) using modules.
- **main.tf, variables.tf, outputs.tf:** Use these in each environment for clarity and organization.
- **Version Control Modules:** Reference modules using Git tags or paths to ensure consistency.





**Explaining how modules
enhance reusability and
maintainability in
Terraform.**

Let's see

Modules in Terraform enhance:

- **Reusability:** Write once, use across multiple environments or projects (e.g., a VPC module).
- **Maintainability:** Centralize updates—changes in a module apply wherever it's used.
- **Organization:** Break infrastructure into logical, manageable components.
- **Consistency:** Enforce standards and reduce configuration drift.





**Discussing how
workspaces separate
environments (dev,
staging, prod) for
infrastructure.**

Let's discuss

Terraform workspaces help separate environments like dev, staging, and prod by:

- Maintaining separate state files for each workspace, isolating resources.
- Using the same codebase to manage different environments.
- Enabling safe testing and deployment without affecting production.





**Demonstrating how to
use Terraform modules
for AWS networking,
compute, and storage.**

Let's do it

To use Terraform modules for AWS networking, compute, and storage:

- **Networking Module:**

```
module "vpc" {  
  source = "./modules/vpc"  
  cidr_block = "10.0.0.0/16"  
}
```



Let's do it

- **Compute Module:**

```
module "ec2" {  
    source = "./modules/ec2"  
    instance_type = "t2.micro"  
    ami_id = "ami-123456"  
}
```

- **Storage Module:**

```
module "s3" {  
    source = "./modules/s3"  
    bucket_name = "my-app-bucket"  
}
```





**Explaining why storing
credentials in Terraform
code is a bad practice.**

Let's see

Storing credentials in Terraform code is a bad practice because:

- Security Risk
- Version Control Issues
- Lack of Flexibility





**Demonstrating using
environment variables
and HashiCorp Vault for
secrets management.**

Let's do it

To use environment variables and HashiCorp Vault for secrets management in Terraform:

1. Environment Variables:

- Set environment variables:

```
export AWS_ACCESS_KEY_ID="your-access-key"  
export AWS_SECRET_ACCESS_KEY="your-secret-key"
```

- Reference in Terraform:

```
provider "aws" {  
    access_key = var.AWS_ACCESS_KEY_ID  
    secret_key = var.AWS_SECRET_ACCESS_KEY  
}
```



Let's do it

2. HashiCorp Vault:

- Authenticate with Vault:

```
vault login <token>
```

- Retrieve secrets in Terraform:

```
data "vault_secret" "my_secret" {
    path = "secret/myapp"
}

resource "aws_secretsmanager_secret" "example" {
    secret_string = data.vault_secret.my_secret.data["password"]
}
```





**Discussing role-based
access control (RBAC) for
Terraform workflows.**

Let's discuss

Role-Based Access Control (RBAC) for Terraform workflows:

- Defines Permissions
- Least Privilege
- Separation of Duties
- Integration with Cloud Providers



Pop Quiz

Q. What is the primary purpose of RBAC in Terraform workflows?

A

To manage networking and
compute resources

B

To define granular permissions for
users and teams

Pop Quiz

Q. What is the primary purpose of RBAC in Terraform workflows?

A

To manage networking and
compute resources

B

To define granular permissions for
users and teams



Take A 5-Minute Break!



- Stretch and relax
- Hydrate
- Clear your mind
- Be back in 5 minutes





**Explaining the risks of
storing Terraform state
locally and why
encryption is necessary.**

Let's see

Risks of storing Terraform state locally:

- **Data Loss:** Local files can be accidentally deleted or lost.
- **Collaboration Issues:** Teams can't safely share or manage local state.
- **Security Risk:** State files often contain sensitive data (e.g., passwords, resource metadata).

Why encryption is necessary:

- **Protects Sensitive Info:** Encrypts secrets stored in the state file.
- **Prevents Unauthorized Access:** Secures data at rest, especially in remote backends like S3.
- **Compliance:** Helps meet security and regulatory requirements.



**Demonstrating storing
Terraform state in AWS
S3 with encryption
enabled.**

Let's do it

To store Terraform state in AWS S3 with encryption enabled:

```
terraform {  
  backend "s3" {  
    bucket      = "my-terraform-state-bucket"  
    key         = "envs/prod/terraform.tfstate"  
    region      = "us-east-1"  
    encrypt     = true  
    dynamodb_table = "terraform-locks" # for state locking  
  }  
}
```

- `encrypt = true`: Enables server-side encryption.
- Use dynamodb table for locking to avoid conflicts in team environments.





**Discussing using IAM
roles/policies to restrict
state file access to
authorized users.**

Let's discuss

Using IAM roles/policies restricts access to Terraform state files by:

- Controlling Access
- Enforcing Least Privilege
- Audit & Compliance

Example policy snippet:

```
{  
  "Effect": "Allow",  
  "Action": ["s3:GetObject", "s3:PutObject"],  
  "Resource": "arn:aws:s3:::my-terraform-state-bucket/*"  
}
```





**Explaining how
Terraform interacts with
AWS using the AWS
provider.**

Let's see

Terraform interacts with AWS using the AWS provider by:

- Authenticating via AWS credentials (env vars, IAM roles, or config files).
- Mapping Terraform code to AWS resources, like EC2, S3, VPC, etc.
- Using the AWS API to create, update, or delete resources based on the declared infrastructure.

Example:

```
provider "aws" {  
    region = "us-east-1"  
}
```



Pop Quiz

Q. What is the purpose of the AWS provider in Terraform?

A

To manage Terraform state files

B

To allow Terraform to interact with
AWS services

Pop Quiz

Q. What is the purpose of the AWS provider in Terraform?

A

To manage Terraform state files

B

To allow Terraform to interact with
AWS services



**Demonstrating writing
Terraform configurations
to provision (Provide
pre-built configurations)**

Let's do it

1. EC2 with Security Group:

```
resource "aws_security_group" "web_sg" {
  name          = "web-sg"
  description   = "Allow HTTP and SSH"
  ingress {
    from_port   = 80
    to_port     = 80
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

resource "aws_instance" "web" {
  ami           = "ami-123456"
  instance_type = "t2.micro"
  security_groups = [aws_security_group.web_sg.name]
}
```



Let's do it

2. S3 with Versioning & Encryption:

```
resource "aws_s3_bucket" "my_bucket" {
  bucket = "my-secure-bucket"
}

resource "aws_s3_bucket_versioning" "versioning" {
  bucket = aws_s3_bucket.my_bucket.id
  versioning_configuration {
    status = "Enabled"
  }
}

resource "aws_s3_bucket_server_side_encryption_configuration" "encryption" {
  bucket = aws_s3_bucket.my_bucket.id
  rule {
    apply_server_side_encryption_by_default {
      sse_algorithm = "AES256"
    }
  }
}
```



Let's do it

3. VPC with Subnets & Route Tables:

```
resource "aws_vpc" "main" {
    cidr_block = "10.0.0.0/16"
}

resource "aws_subnet" "subnet1" {
    vpc_id      = aws_vpc.main.id
    cidr_block = "10.0.1.0/24"
}

resource "aws_route_table" "rt" {
    vpc_id = aws_vpc.main.id
}

resource "aws_route_table_association" "rta" {
    subnet_id      = aws_subnet.subnet1.id
    route_table_id = aws_route_table.rt.id
}
```



Let's do it

4. RDS with Automated Backups:

```
resource "aws_db_instance" "db" {  
    identifier      = "mydb"  
    engine          = "mysql"  
    instance_class = "db.t3.micro"  
    allocated_storage = 20  
    username        = "admin"  
    password        = "password123"  
    skip_final_snapshot = true  
    backup_retention_period = 7  
    backup_window           = "03:00-06:00"  
}
```





Explaining Auto Scaling Groups (ASG) and Elastic Load Balancing (ELB).

Let's see

Auto Scaling Groups (ASG):

- Automatically adjust the number of EC2 instances based on demand to ensure performance and cost-efficiency.

Elastic Load Balancing (ELB):

- Distributes incoming traffic across multiple EC2 instances to improve availability, fault tolerance, and scalability.





**Demonstrating how
Terraform can
dynamically scale EC2
instances based on CPU
usage.**

Let's do it

To dynamically scale EC2 instances based on CPU usage with Terraform:

```
resource "aws_autoscaling_group" "web_asg" {
  launch_configuration = aws_launch_configuration.web_lc.name
  min_size              = 1
  max_size              = 5
  desired_capacity      = 2
  target_group_arns     = [aws_lb_target_group.web_tg.arn]
  vpc_zone_identifier   = ["subnet-abc123"]
}

resource "aws_autoscaling_policy" "cpu_scale_up" {
  name                  = "scale-up"
  scaling_adjustment    = 1
  adjustment_type       = "ChangeInCapacity"
  cooldown              = 300
  autoscaling_group_name = aws_autoscaling_group.web_asg.name
}

resource "aws_cloudwatch_metric_alarm" "high_cpu" {
  alarm_name            = "high-cpu"
  comparison_operator   = "GreaterThanOrEqualToThreshold"
  evaluation_periods    = 2
  metric_name           = "CPUUtilization"
  namespace              = "AWS/EC2"
  period                 = 120
  statistic              = "Average"
  threshold              = 70
  alarm_actions          = [aws_autoscaling_policy.cpu_scale_up.arn]
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.web_asg.name
  }
}
```



**Discussing best practices
for high availability and
failover using ASG and
ELB.**

Let's discuss

Best practices for high availability and failover with ASG & ELB:

- Use Multi-AZ Deployment
- Health Checks
- Auto Scaling
- Elastic Load Balancer
- Monitoring & Alerts



Pop Quiz

Q. How does an Elastic Load Balancer (ELB) contribute to high availability?

A

It distributes incoming traffic across
multiple healthy targets

B

It replaces unhealthy instances
automatically

Pop Quiz

Q. How does an Elastic Load Balancer (ELB) contribute to high availability?

A

It distributes incoming traffic across multiple healthy targets

B

It replaces unhealthy instances automatically



Explaining the importance of monitoring and logging in cloud environments.

Let's see

Monitoring and logging in cloud environments are crucial for:

- Visibility
- Troubleshooting
- Security
- Compliance
- Optimization





**Discussing how
Terraform can automate
the setup of CloudWatch
dashboards, alarms, and
log groups.**

Let's discuss

Terraform can automate CloudWatch setup by:

- **Dashboards:** Create visual summaries of metrics.

```
resource "aws_cloudwatch_dashboard" "main" {  
    dashboard_name = "app-dashboard"  
    dashboard_body = jsonencode({ ... })  
}
```



Let's discuss

- **Alarms:** Trigger actions based on thresholds (e.g., high CPU).

```
resource "aws_cloudwatch_metric_alarm" "cpu_alarm" { ... }
```

- **Log Groups:** Collect and manage logs for resources.

```
resource "aws_cloudwatch_log_group" "app_logs" {  
    name          = "/aws/app/logs"  
    retention_in_days = 7  
}
```





**Demonstrating
configuring CloudWatch
alarms for instance
health checks and
resource utilization.**

Let's do it

To configure CloudWatch alarms for instance health and resource usage:

```
resource "aws_cloudwatch_metric_alarm" "ec2_cpu_high" {
    alarm_name          = "HighCPUUsage"
    metric_name         = "CPUUtilization"
    namespace           = "AWS/EC2"
    statistic            = "Average"
    period              = 300
    evaluation_periods = 2
    threshold           = 80
    comparison_operator = "GreaterThanThreshold"
    dimensions = {
        InstanceId = "i-0123456789abcdef0"
    }
    alarm_actions = ["arn:aws:sns:region:account-id:alarm-topic"]
}
```



Explaining how EBS snapshots work for disaster recovery.

Let's see

EBS snapshots help in disaster recovery by:

- **Backing up volumes:** Capture point-in-time copies of EBS volumes.
- **Storing in S3:** Snapshots are stored durably and securely in Amazon S3.
- **Restoring quickly:** Use snapshots to create new volumes in any AZ.
- **Automation:** Schedule regular snapshots for consistent backups.





**Demonstrating using
Terraform to schedule
automated EBS
snapshots.**

Let's do it

To schedule automated EBS snapshots with Terraform, use Data Lifecycle Manager (DLM):

```
resource "aws_dlm_lifecycle_policy" "ebs_backup" {
  description = "Daily snapshot of EBS volumes"
  execution_role_arn = "arn:aws:iam::123456789012:role/AWSDataLifecycleManagerDefaultRole"

  policy_details {
    resource_types = ["VOLUME"]
    target_tags = {
      Backup = "true"
    }
    schedules {
      name = "daily-backup"
      create_rule {
        interval      = 24
        interval_unit = "HOURS"
      }
      retain_rule {
        count = 7
      }
    }
  }
}
```



**Discuss failover
mechanisms using
multi-region S3
replication and RDS read
replicas.**

Let's discuss

Failover mechanisms using:

S3 Multi-Region Replication:

- Automatically replicates objects across AWS regions for durability and quick recovery if a region fails.

RDS Read Replicas:

- Create read-only copies in different regions for load balancing and promote them to primary in case of failure, ensuring minimal downtime.





Time for case study!

Important

- Complete the post-class assessment
- Complete assignments (if any)
- Practice the concepts and techniques taught in this session
- Review your lecture notes
- Note down questions and queries regarding this session and consult the teaching assistants



Thanks



!

